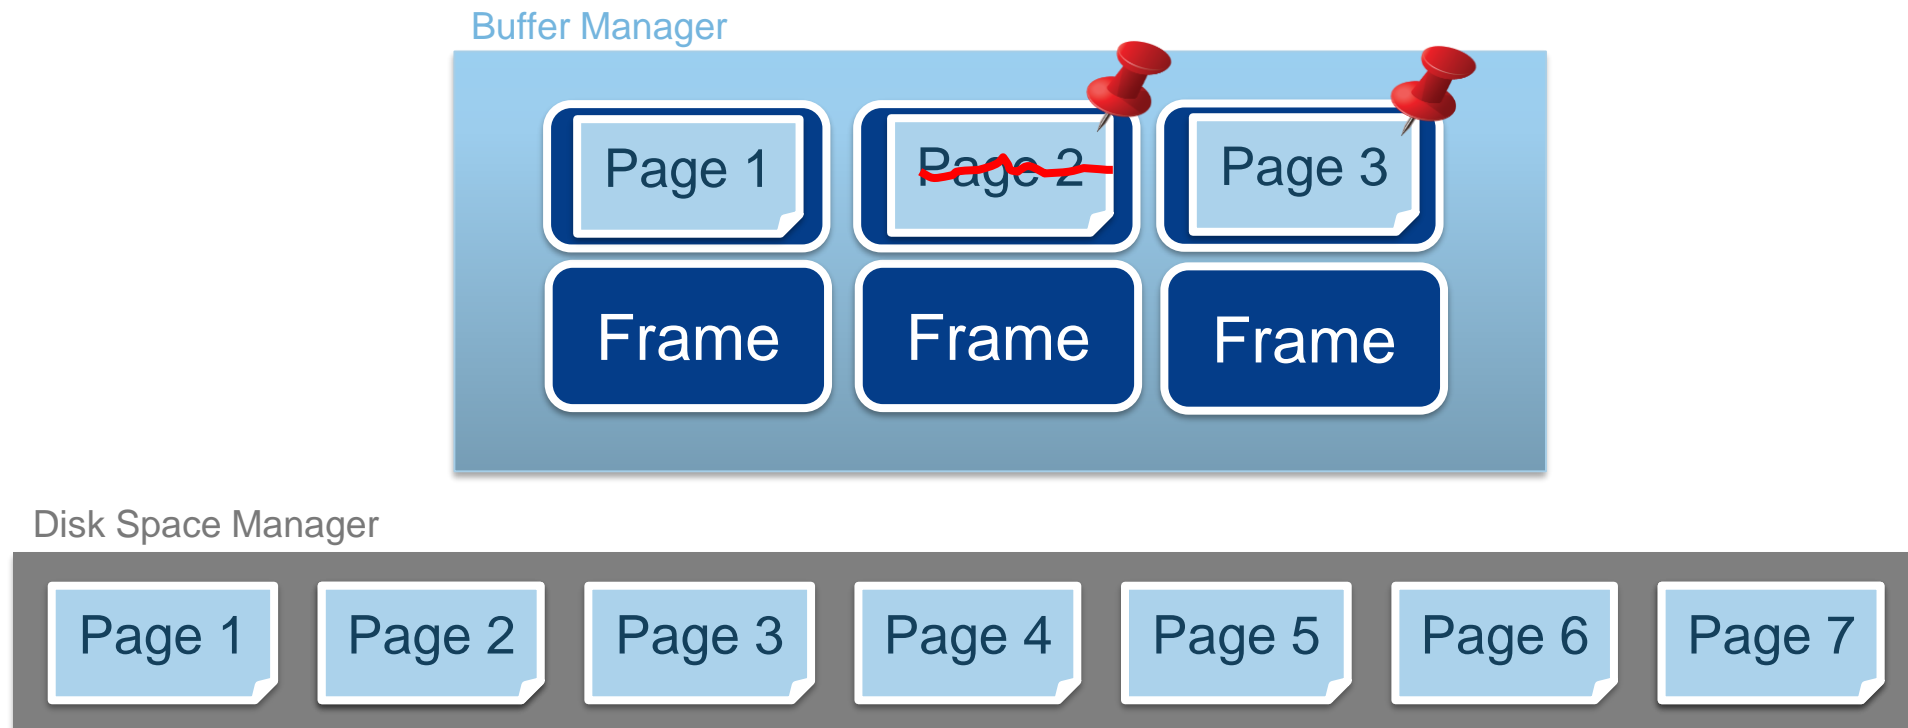


Project3

Buffer management

Buffer Management

- Current disk-based b+tree doesn't support buffer management.
- Our goal is to implement **in-memory buffer manager** to caching on-disk pages.



Project Specification

- Define the buffer block structure, which must contain at least those fields.
- **Physical frame:** containing up to date contents of target page.
 - **Table id:** the unique id of table (per file)
 - **Page number:** the target page number within a file.
 - **Is dirty:** whether this buffer block is dirty or not.
 - **Is pinned:** whether this buffer is accessed right now.
 - **LRU list next (prev) :** buffer blocks are managed by LRU list.
 - Other information can be added with your own buffer manager design.

Buffer Structure

frame (page size : 4096 bytes)
table_id
page_num
is_dirty
is_pinned
next/prev of LRU

Project Specification

➤ Implement database initialization function.

- **int init_db (int num_buf);**
- Allocate the buffer pool (array) with the given number of entries.
- Initialize other fields (state info, LRU info..) with your own design.
- If success, return 0. Otherwise, return non-zero value.

➤ **open_table** interface

- **int open_table (char *pathname);**
- Open existing data file or create one if not existed. You can give same table id when db opens same table more than once after init_db.
- If success, return the **unique table id**, which represents the own table in this database. (Return negative value if error occurs)
- You have to maintain a table id once open_table() is called, which is matching file descriptor or file pointer depending on your previous implementation. (table id ≥ 1 and maximum allocated id is set to 10)

Project Specification

- Table id is also used in previous **insert**, **delete**, **find** interfaces as well. Modify those to table APIs.
 - `int db_insert (int table_id, int64_t key, char * value);`
 - `int db_find (int table_id, int64_t key, char * ret_val);`
 - `int db_delete (int table_id, int64_t key);`
- Your existing APIs (insert, delete, find) must work with implemented buffer manager first before accessing to disk. (more details in next slides)
 - If the page is not in buffer pool (cache-miss), read page from disk and maintain this page in buffer block.
 - Page modification only occurs in-memory buffer. If the page frame in buffer is updated, mark the buffer block as dirty.
 - According to LRU policy, least recently used buffer is the victim for page eviction. Writing page to disk occurs during LRU page eviction.

Project Specification

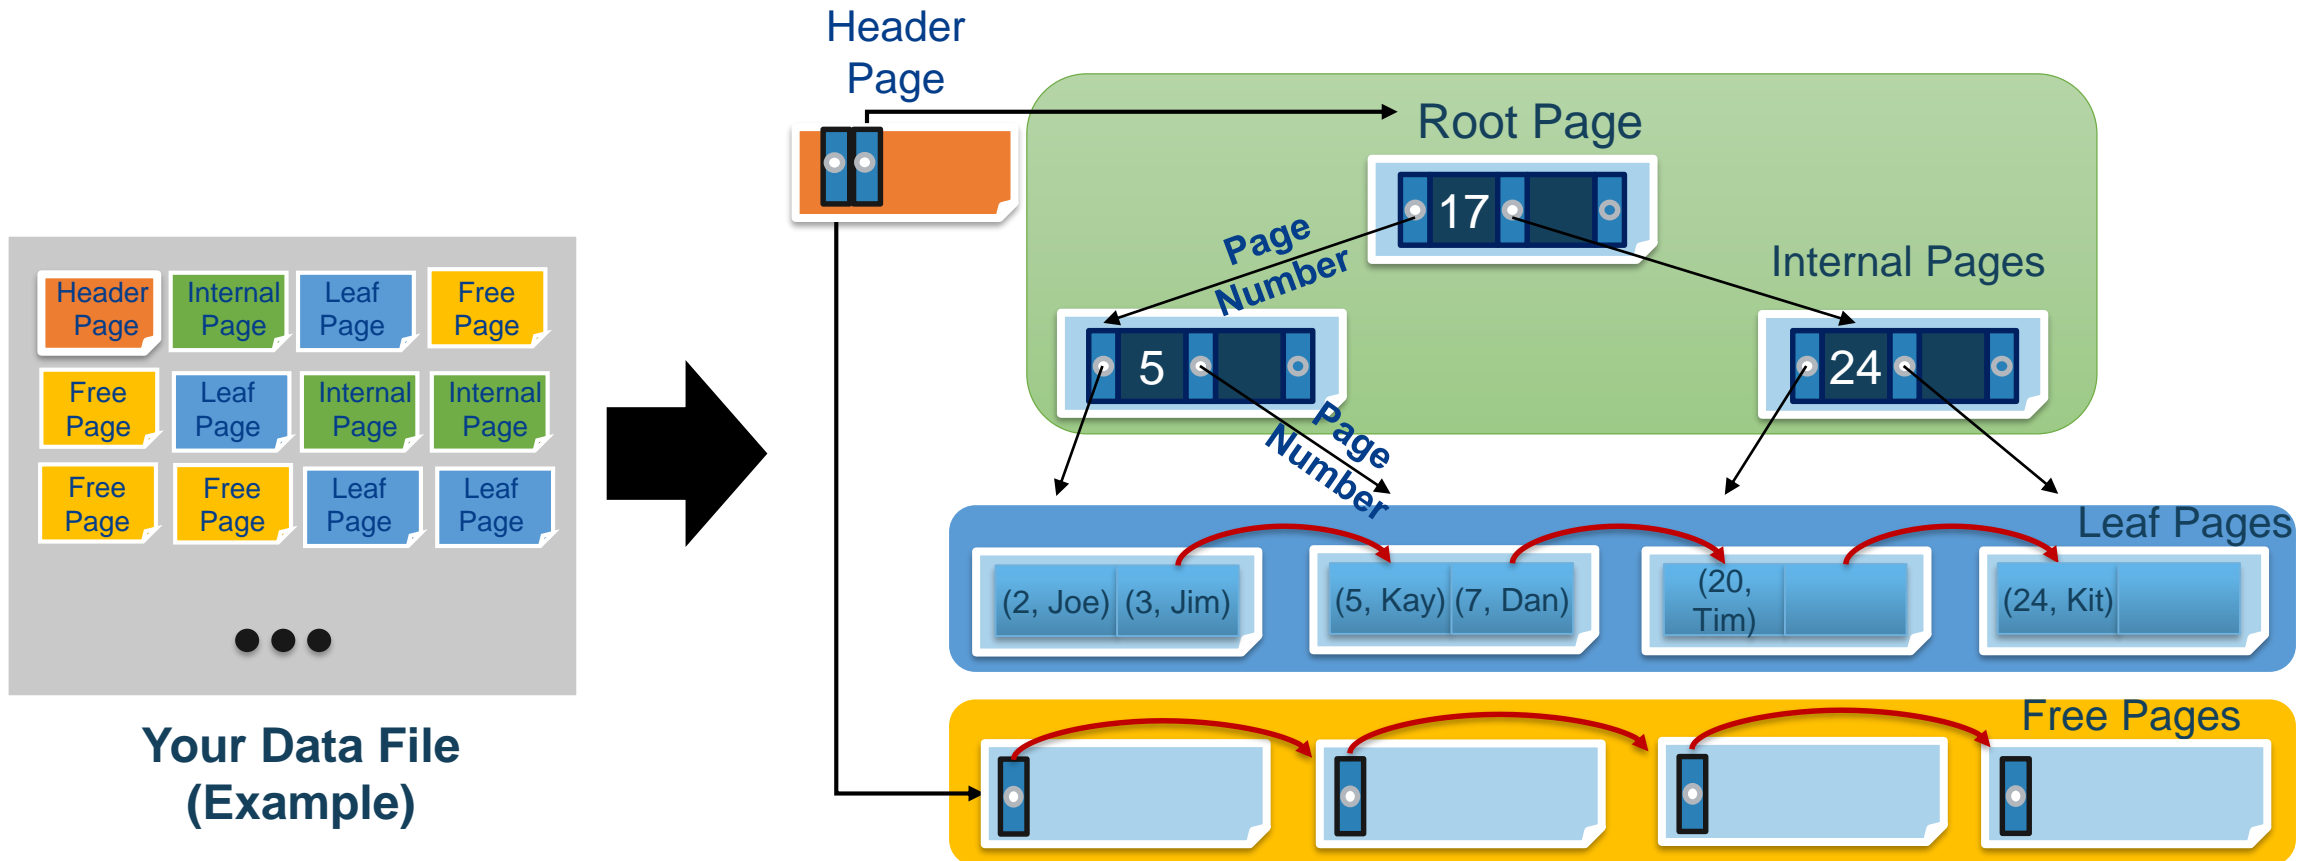
- Implement **close_table** interface.
 - **int close_table(int table_id);**
 - Write all pages of this table from buffer to disk and discard the table id.
 - If success, return 0. Otherwise, return non-zero value.
- Implement database shutdown function.
 - **int shutdown_db();**
 - Flush all data from buffer and destroy allocated buffer.
 - If success, return 0. Otherwise, return non-zero value.

Project Specification

➤ Your library (libbpt.a) should provide those API services.

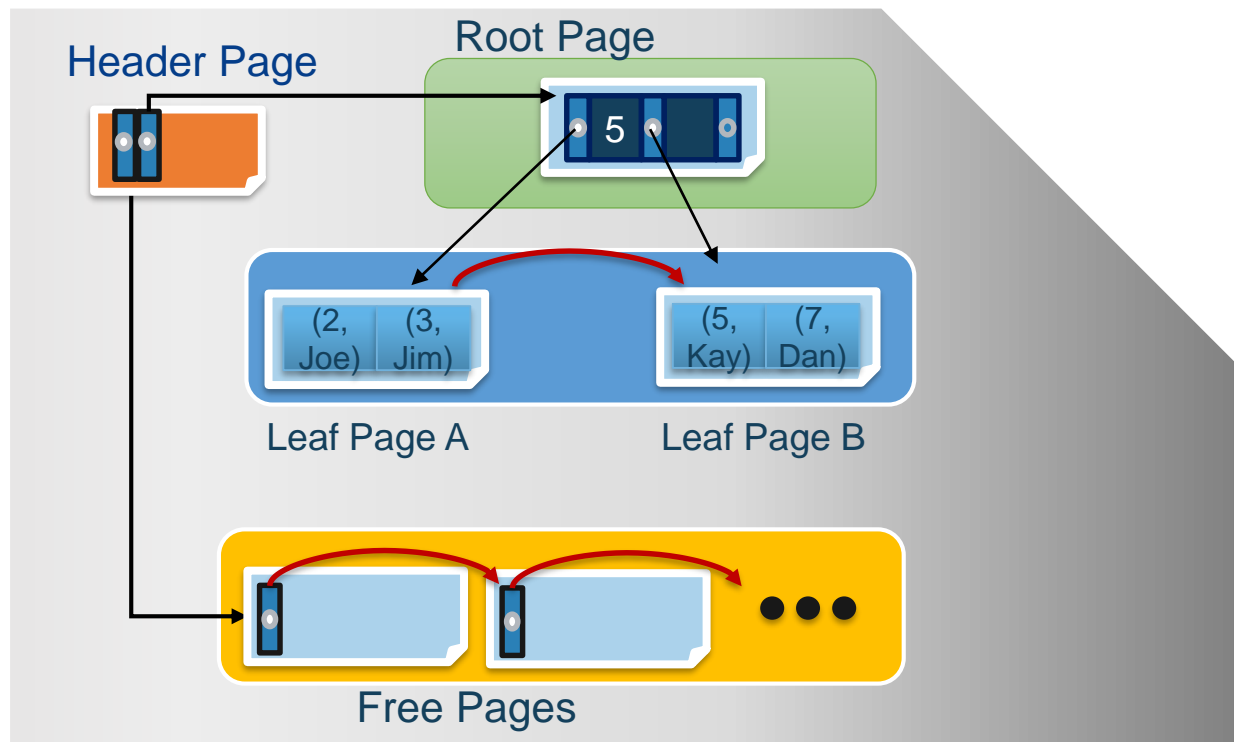
1. **int init_db (int buf_num);**
 - Initialize buffer pool with given number and buffer manager.
2. **int open_table (char * pathname);**
 - Open existing data file using 'pathname' or create one if not existed. If success, return **table_id**.
3. **int db_insert (int table_id, int64_t key, char * value);**
4. **int db_find (int table_id, int64_t key, char * ret_val);**
5. **int db_delete (int table_id, int64_t key);**
6. **int close_table(int table_id);**
 - Write the pages relating to this table to disk and close the table.
7. **int shutdown_db(void);**
 - Destroy buffer manager.

So far..




Buffer Management

- Assume the on-disk pages are stored like below form.




Client

Buffer Manager

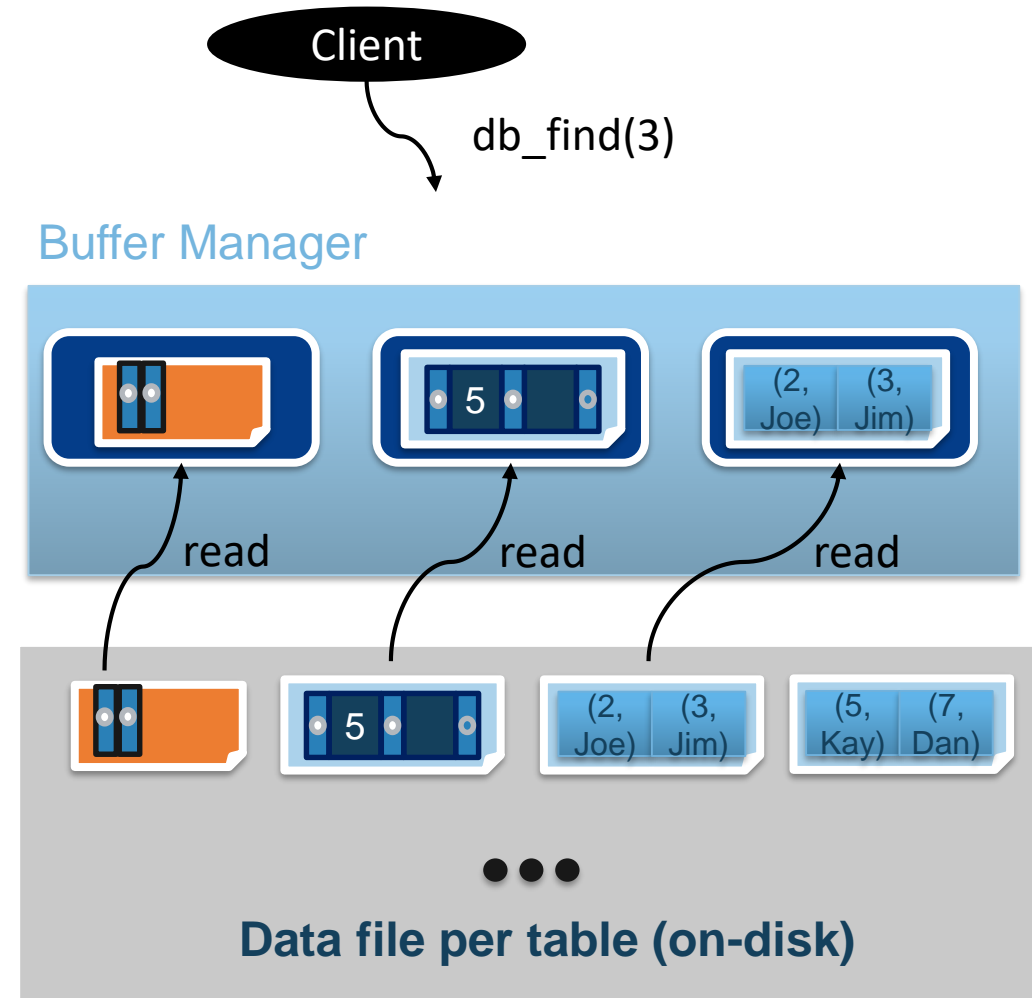


Disk Space Manager



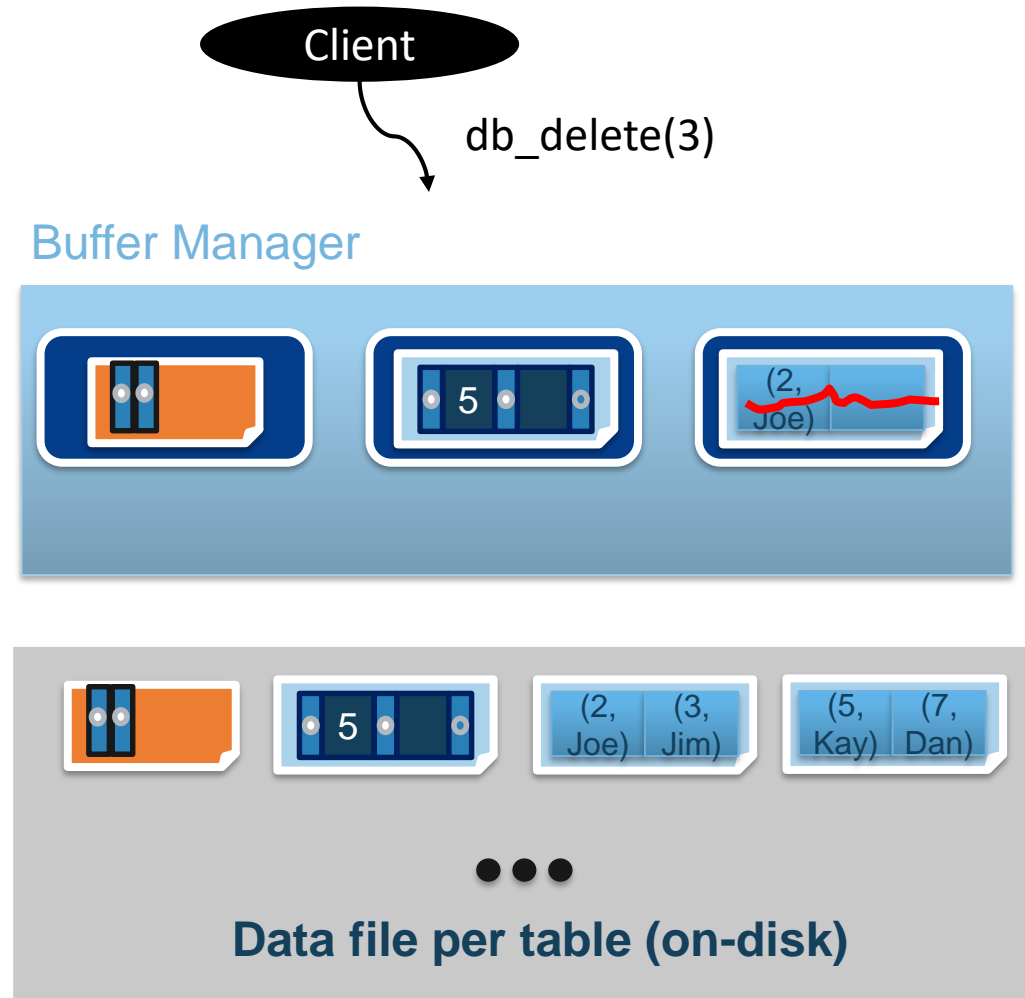
Buffer Management

- First, search the page from the buffer pool.
- If the page is not in the buffer pool (i.e, cache-miss occurs), read the page from disk and maintain this page in buffer block.
- While indexing from root to leaf page A (where key 3 is located), header page and root page (internal page) are also read by buffer manager.



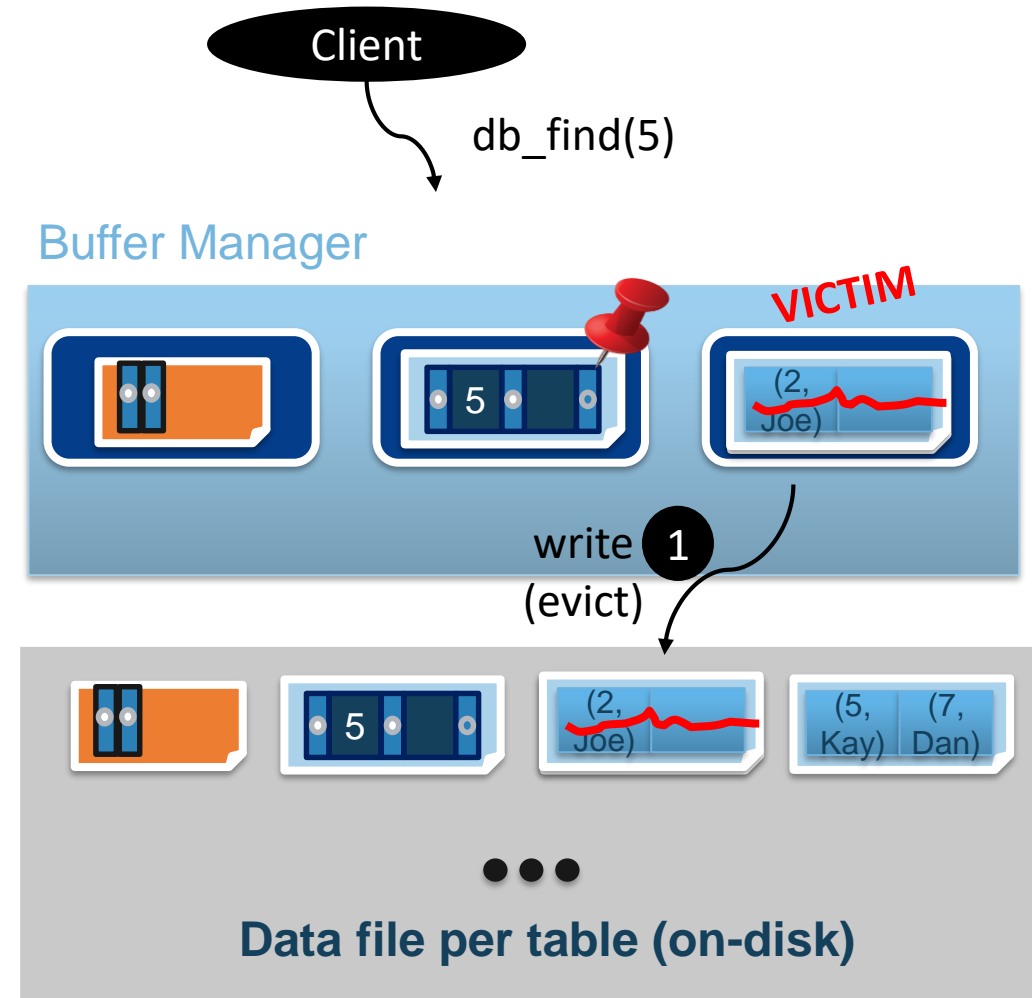
Buffer Management

- After reading page to buffer, update operation can be handled in buffer (memory).
- So delete key 3 operation occurs in buffer, which makes that page marked to dirty.



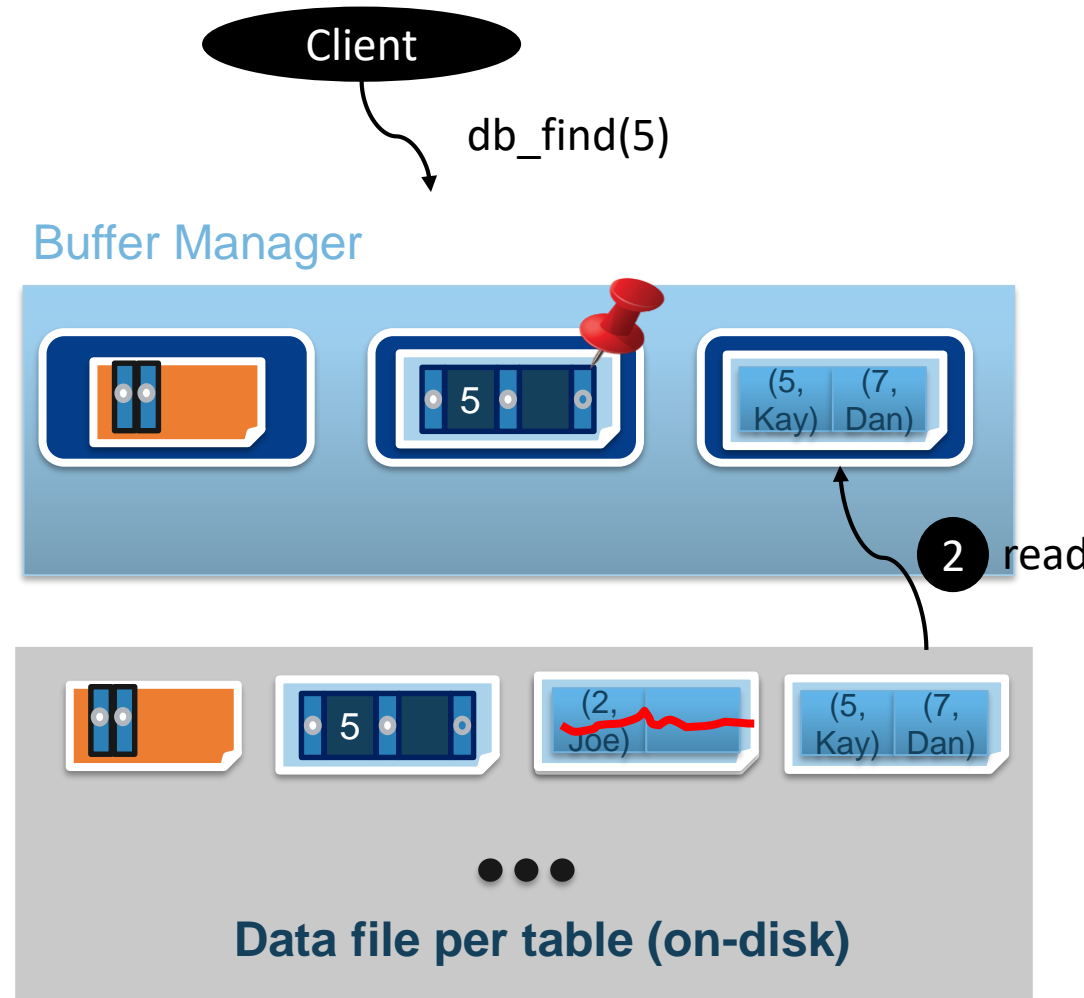
Buffer Management

- Dirty page is written to disk when those page is selected to the victim of LRU policy.
- Assuming example shown left, find(5) tries to read leaf page B which triggers page eviction. (Pinned page should not be the victim of eviction.)
- If the victim page is marked as dirty, write data to disk first. ①



Buffer Management

- Dirty page is written to disk when those page is selected to the victim of LRU policy.
- Assuming example shown left, find(5) tries to read leaf page B which triggers page eviction. (Pinned page should not be the victim of eviction.)
- If the victim page is marked as dirty, write data to disk first. ❶
- Then read another page from disk. ❷



Buffer Management

- `close_table()` or `shutdown_db()` writes out all dirty buffer block to disk.
- `close_table()` writes out the pages only from those relating to given `table_id`.
- This command can provides synchronous semantic (durability) to user, but loses performance.

